

AWS Academy Cloud Architecting

# Module 5: Adding a Database Layer



## Sections

1. Architectural need
2. Database layer considerations
3. Amazon RDS
4. Amazon DynamoDB
5. Database security controls
6. Migrating data into AWS databases

## Demonstration

- Amazon RDS Automated Backup and Read Replicas

## Labs

- Guided Lab: Creating an Amazon RDS Database
- Challenge Lab: Migrating a Database to Amazon RDS



**Knowledge check**

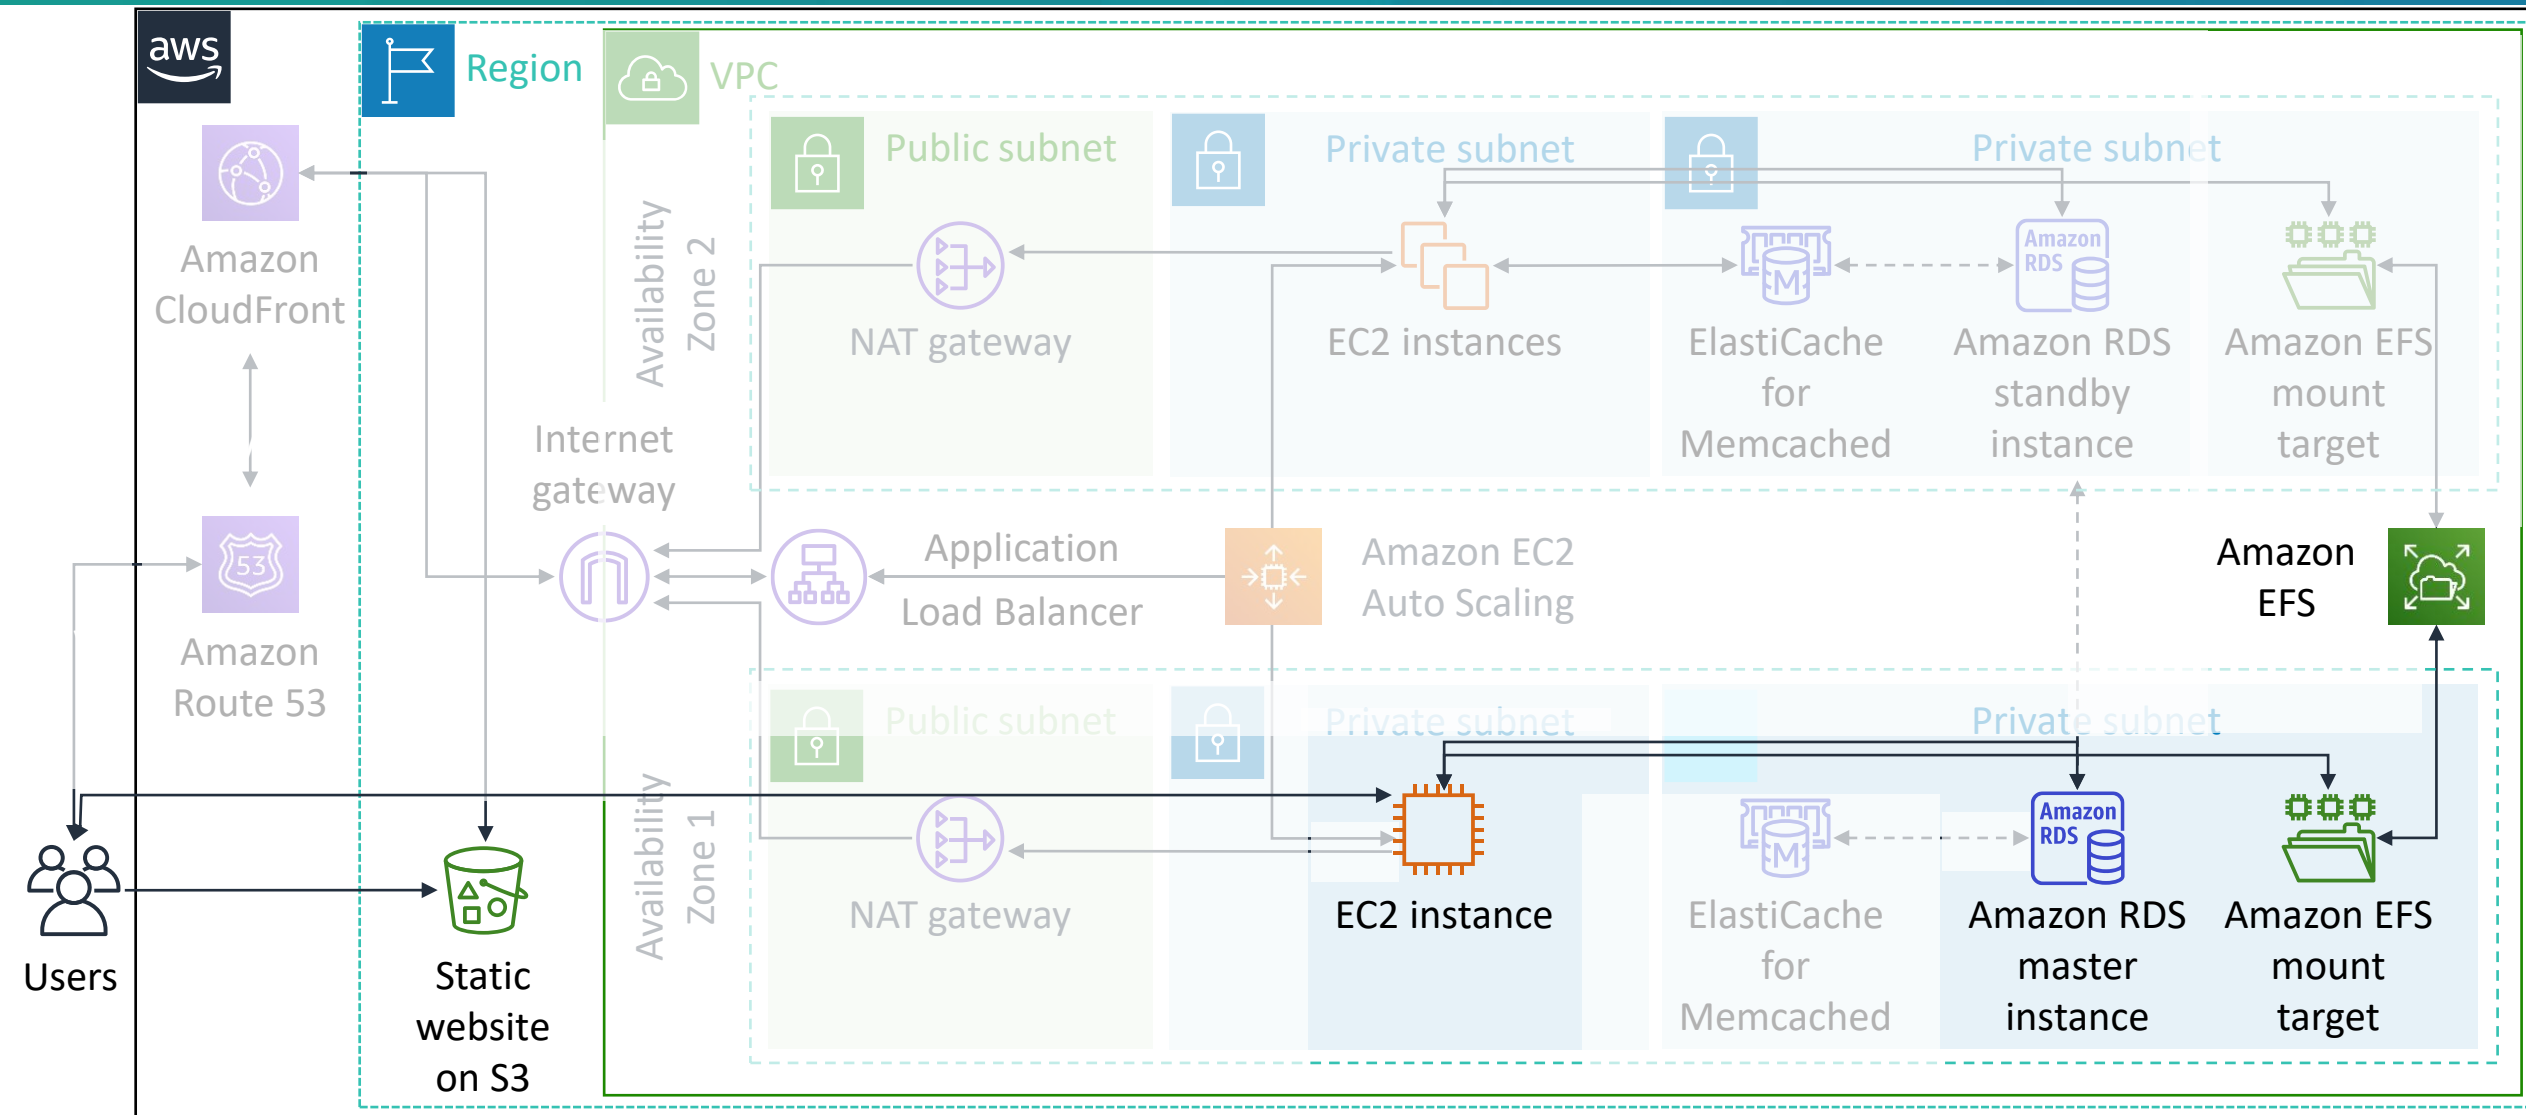
At the end of this module, you should be able to:

- Compare database types
- Differentiate between managed versus unmanaged services
- Explain when to use Amazon Relational Database Service (Amazon RDS)
- Explain when to use Amazon DynamoDB
- Describe available database security controls
- Describe how to migrate data into Amazon Web Services (AWS) databases
- Deploy a database server

## Module 5: Adding a Database Layer

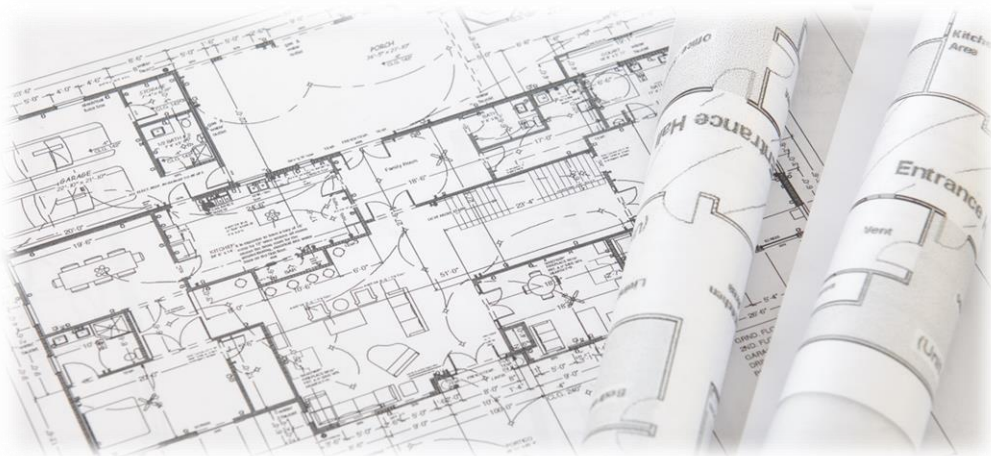
# Section 1: Architectural need

# Databases as part of a larger architecture



# Café business requirement

The café needs a database solution that is easier to maintain, and that provides essential features such as durability, scalability, and high performance.



## Module 5: Adding a Database Layer


# Section 2: Database layer considerations



## Scalability



Total storage requirements



Object size and type



Durability

How much throughput is needed?

Will the chosen solution be able to scale up later, if needed?






# Database considerations: Storage requirements



Scalability



**Total storage requirements**




Object size and type



Durability

**How large does the database need to be?**


**Will it need to store GB, TB, or petabytes of data?**




# Database considerations: Object size and type



Scalability



Total storage requirements




**Object size and type**



Durability


**Do you need to store simple data structures,  
large data objects, or both?**




# Database considerations: Durability



Scalability



Total storage requirements



Object size and type



**Durability**

What level of **data durability**, **data availability**, and **recoverability** is required?

Do regulatory obligations apply?



Now that you reviewed key considerations, consider the two categories of database options available:

## Relational

Traditional examples:

Microsoft SQL Server  
Oracle Database  
MySQL

## Non-Relational

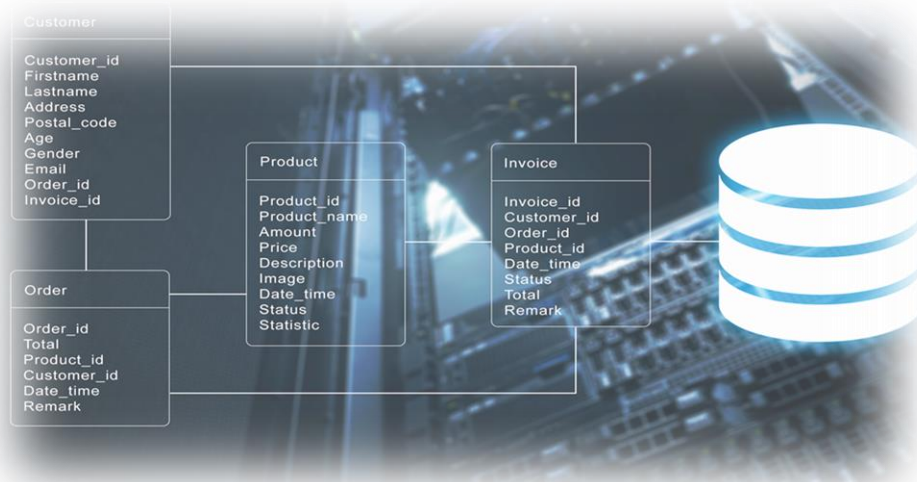
Traditional examples:

MongoDB  
Cassandra  
Redis

# Relational database type

## Benefits:


- Ease of use
- Data integrity
- Reduced data storage
- Common language (structured query language, or SQL)



Relational database management system (RDBMS)

## Relational is ideal when you:

- Need strict schema rules, ACID compliance, and data quality enforcement
- Do not need extreme read/write capacity
- Do not need extreme performance
  - An RDBMS can be the best, lowest-effort solution



# Non-relational database type

## Benefits


- Flexibility
- Scalability
- High performance
- Highly functional APIs

## Non-relational is ideal when:


- Database must scale **horizontally** to handle massive data volume
- Data does not lend itself well to traditional schemas
- Read/write rates exceed what can be economically supported through traditional RDBMS

## Example models

### key-value



### graph




# Amazon database options

*More database options exist—these options are common examples*

## Relational databases



Amazon  
RDS




Amazon  
Redshift



Amazon  
Aurora

## Non-relational databases



Amazon  
DynamoDB



Amazon  
ElastiCache



Amazon  
Neptune

*Focus in this module*

# Section 2 key takeaways



- When you choose a database, consider scalability, storage requirements, the type and size of objects to be stored, and durability requirements
- Relational databases have strict schema rules, provide data integrity, and support SQL
- Non-relational databases scale horizontally, provide higher scalability and flexibility, and work well for semistructured and unstructured data



## Module 5: Adding a Database Layer

# Section 3: Amazon RDS

Relational

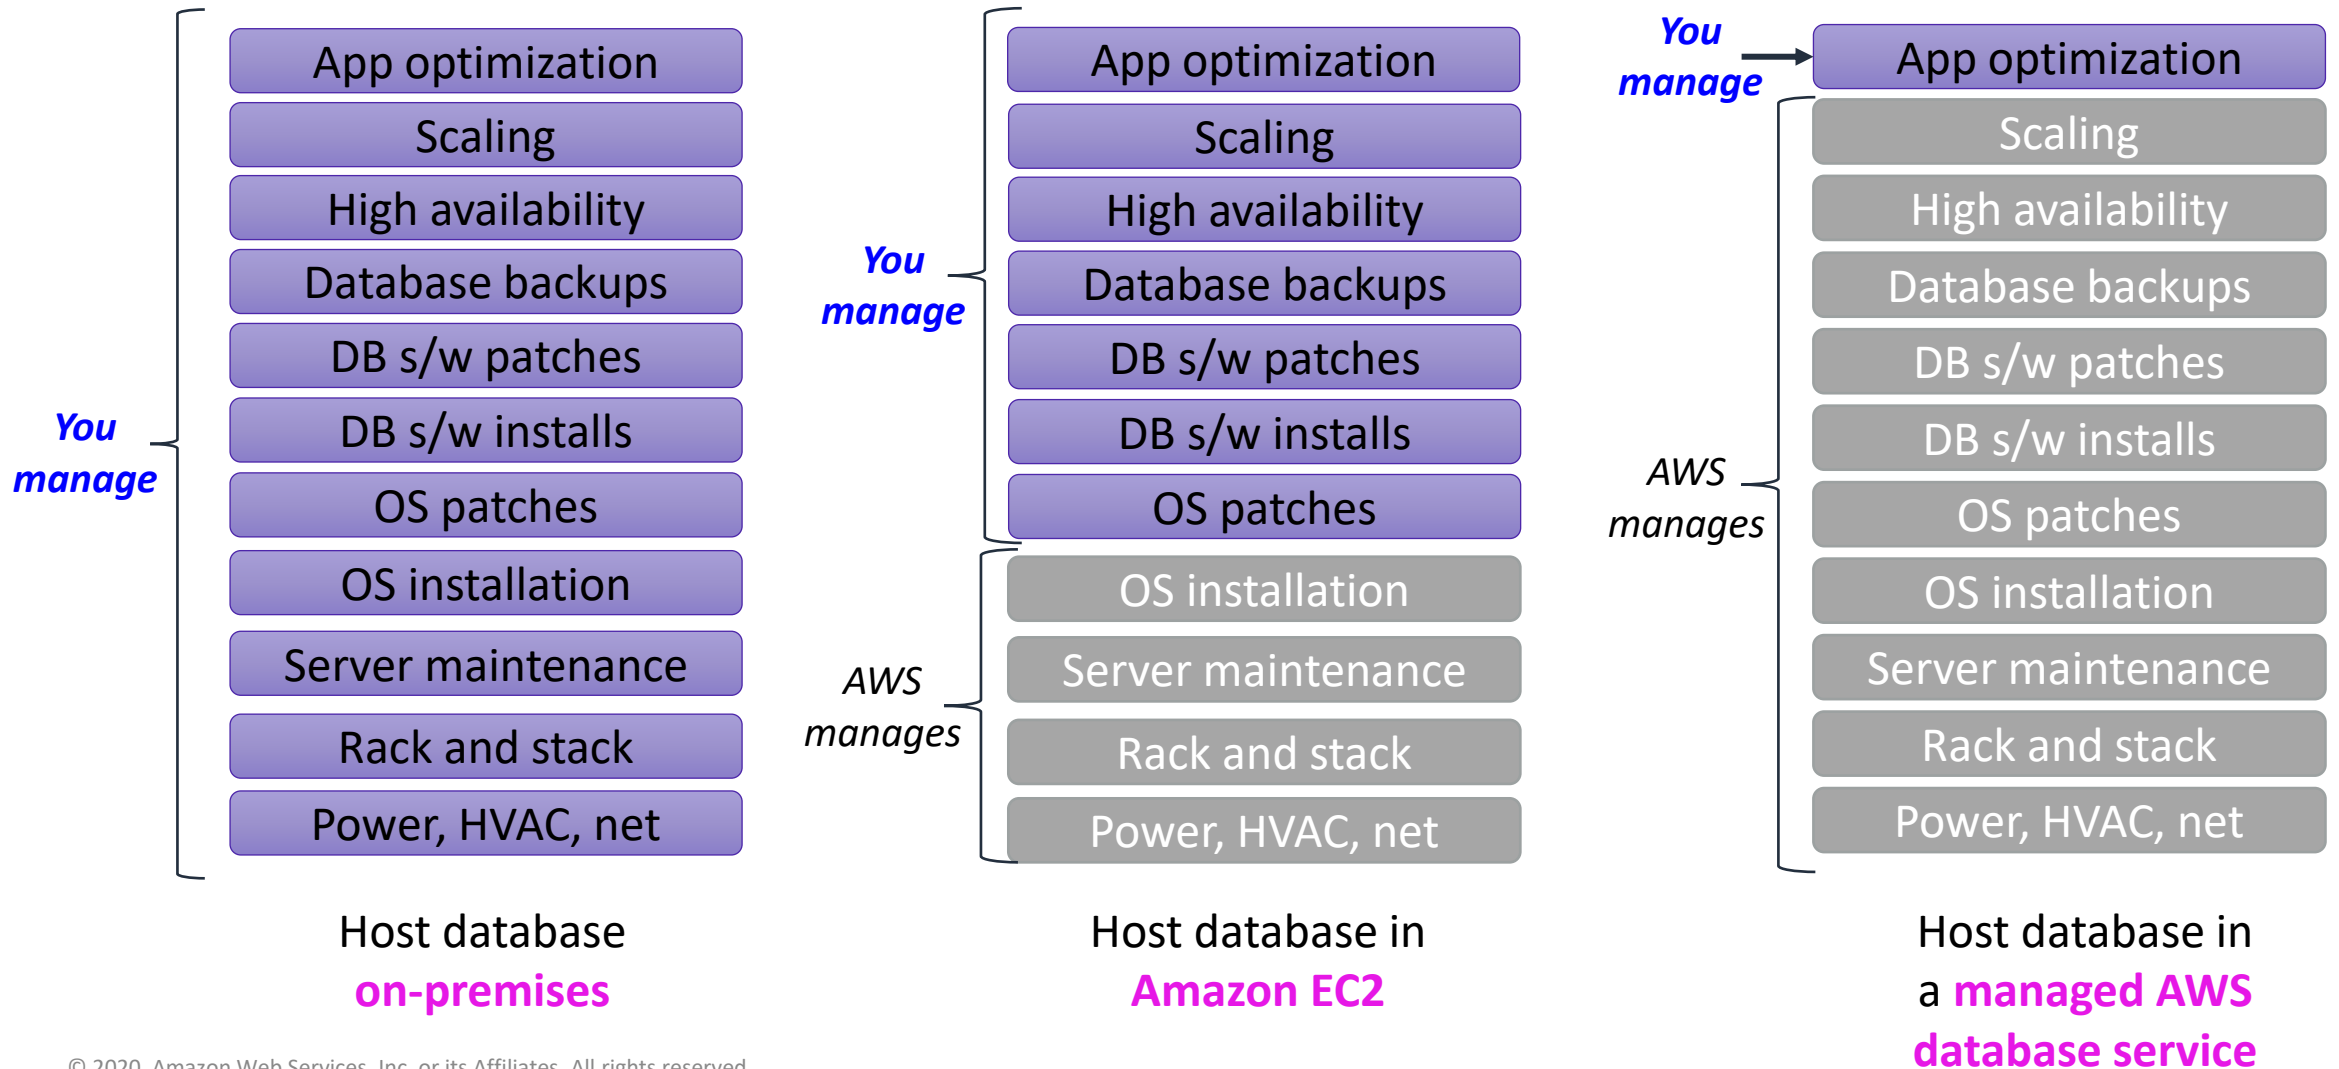


Amazon  
RDS


Amazon RDS is a **fully managed** relational database **service**.



# Advantage of managed AWS database services




# Amazon RDS characteristics



Access pattern


Transactional

Light analytics



Data size

Low-TB range



Performance

Mid to high throughput

Low latency



Business use cases

Transactional

OLAP

# Amazon RDS: Uses and database types



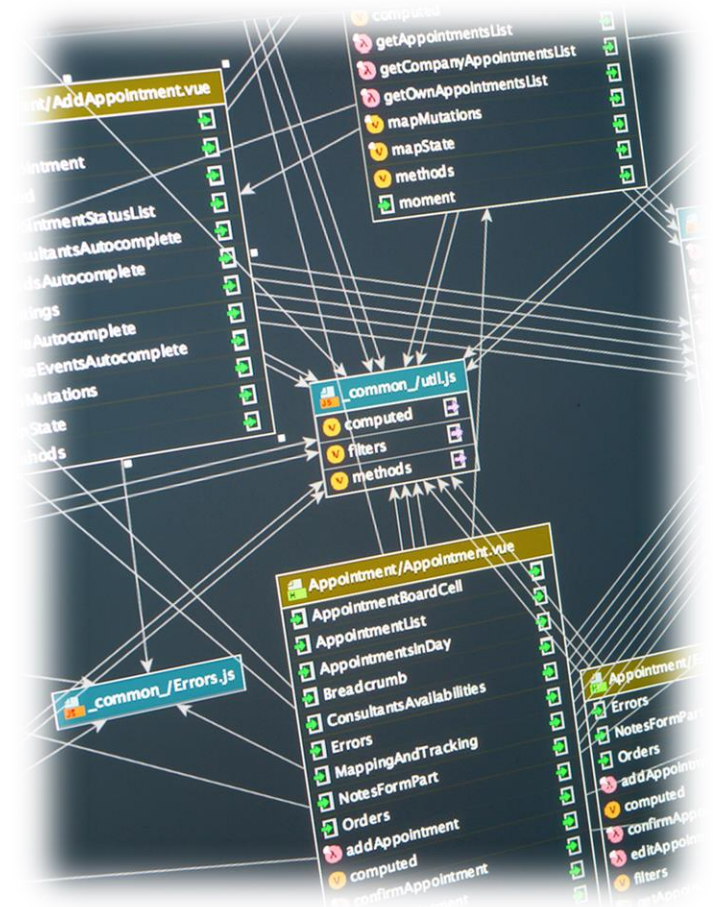
Amazon  
RDS

Works well for applications that:

- Have more complex data
- Need to combine and join datasets
- Need enforced syntax rules

Six database types supported:

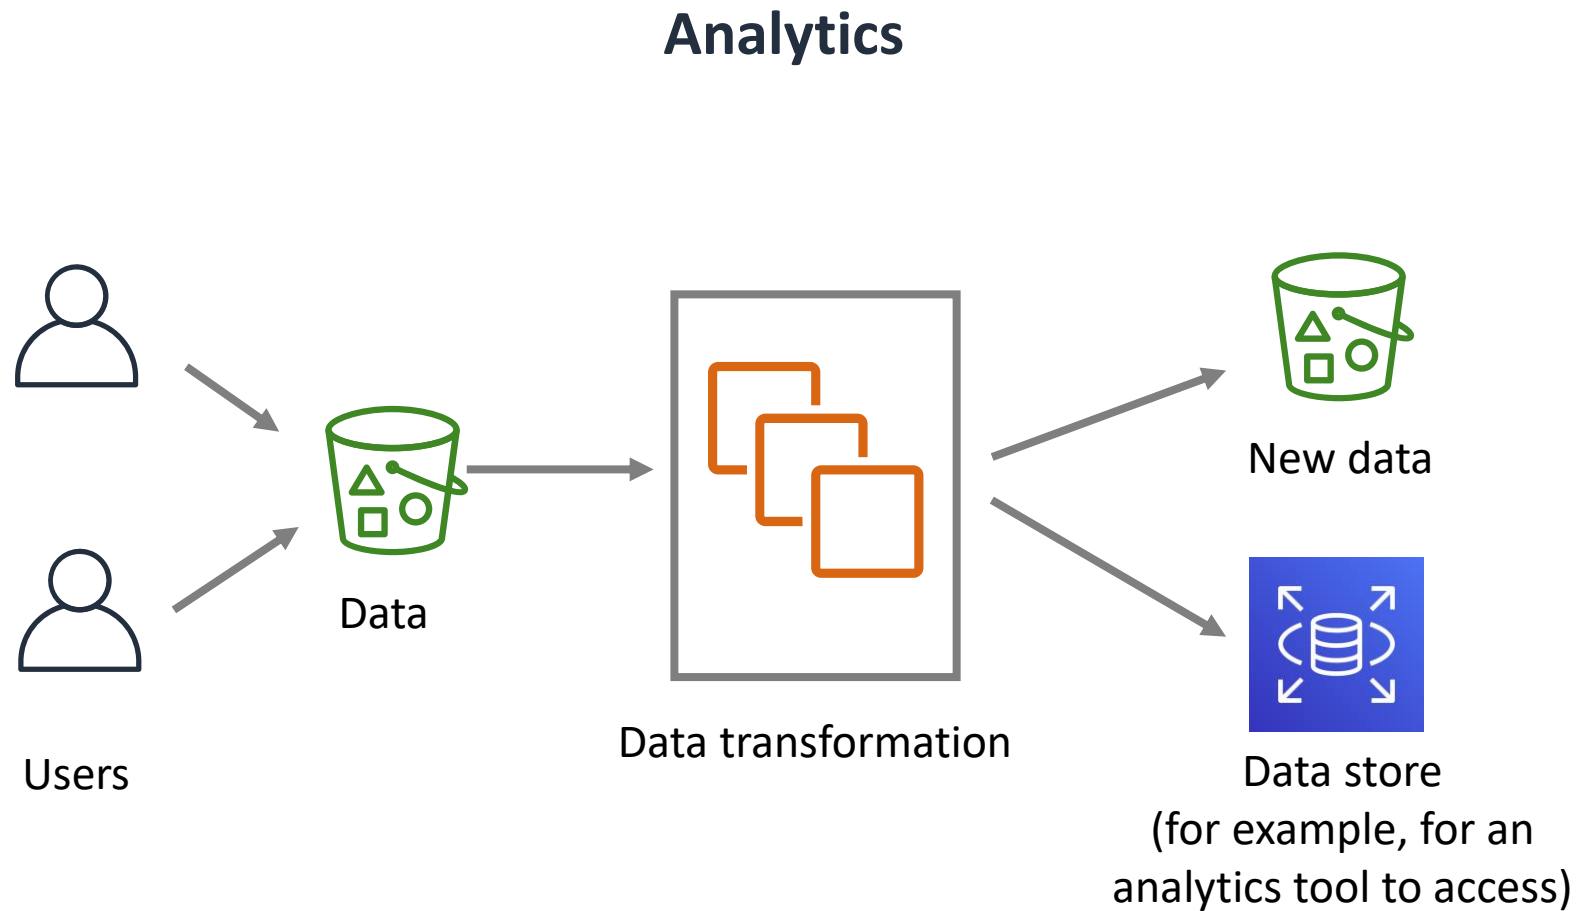
- Microsoft SQL Server
- PostgreSQL
- Oracle
- Aurora
- MySQL
- MariaDB



# Database instance sizing

	T family	M family	R family
Type	Burstable instances	General-purpose instances	Memory-optimized instances
Sizing	1 vCPU/1 GB RAM to 8 vCPU 32 GB RAM	2 vCPU/8 GB RAM to 96 vCPU 384 GB RAM	2 vCPU/16 GB RAM to 96 vCPU 768 GB RAM
Networking	Moderate performance	High performance	High performance
Ideal Workload	Smaller or variable	CPU-intensive	Query-intensive, high connection counts
Highlights	T3 can burst above baseline for extra charge	M5 offers up to 96 vCPU	R5 offers up to 96 vCPU 768 GiB RAM

# Amazon RDS: Example use case





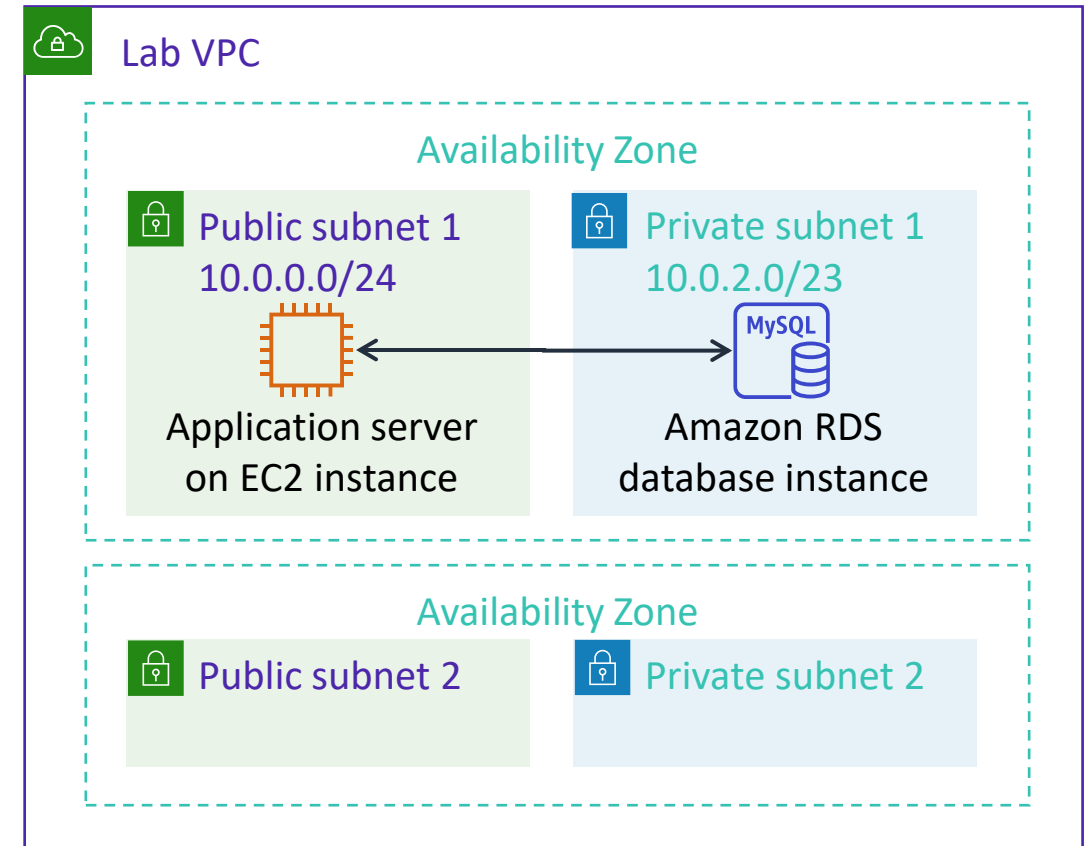
# Module 5 - Guided Lab: Creating an Amazon RDS Database





# Guided lab: Tasks

1. Create an Amazon RDS MySQL database
2. Configure a web application to communicate with the database
3. Modify data in the MySQL database table





# Begin Module 5 – Guided Lab: Creating an Amazon RDS Database



# Guided lab debrief: Key takeaways



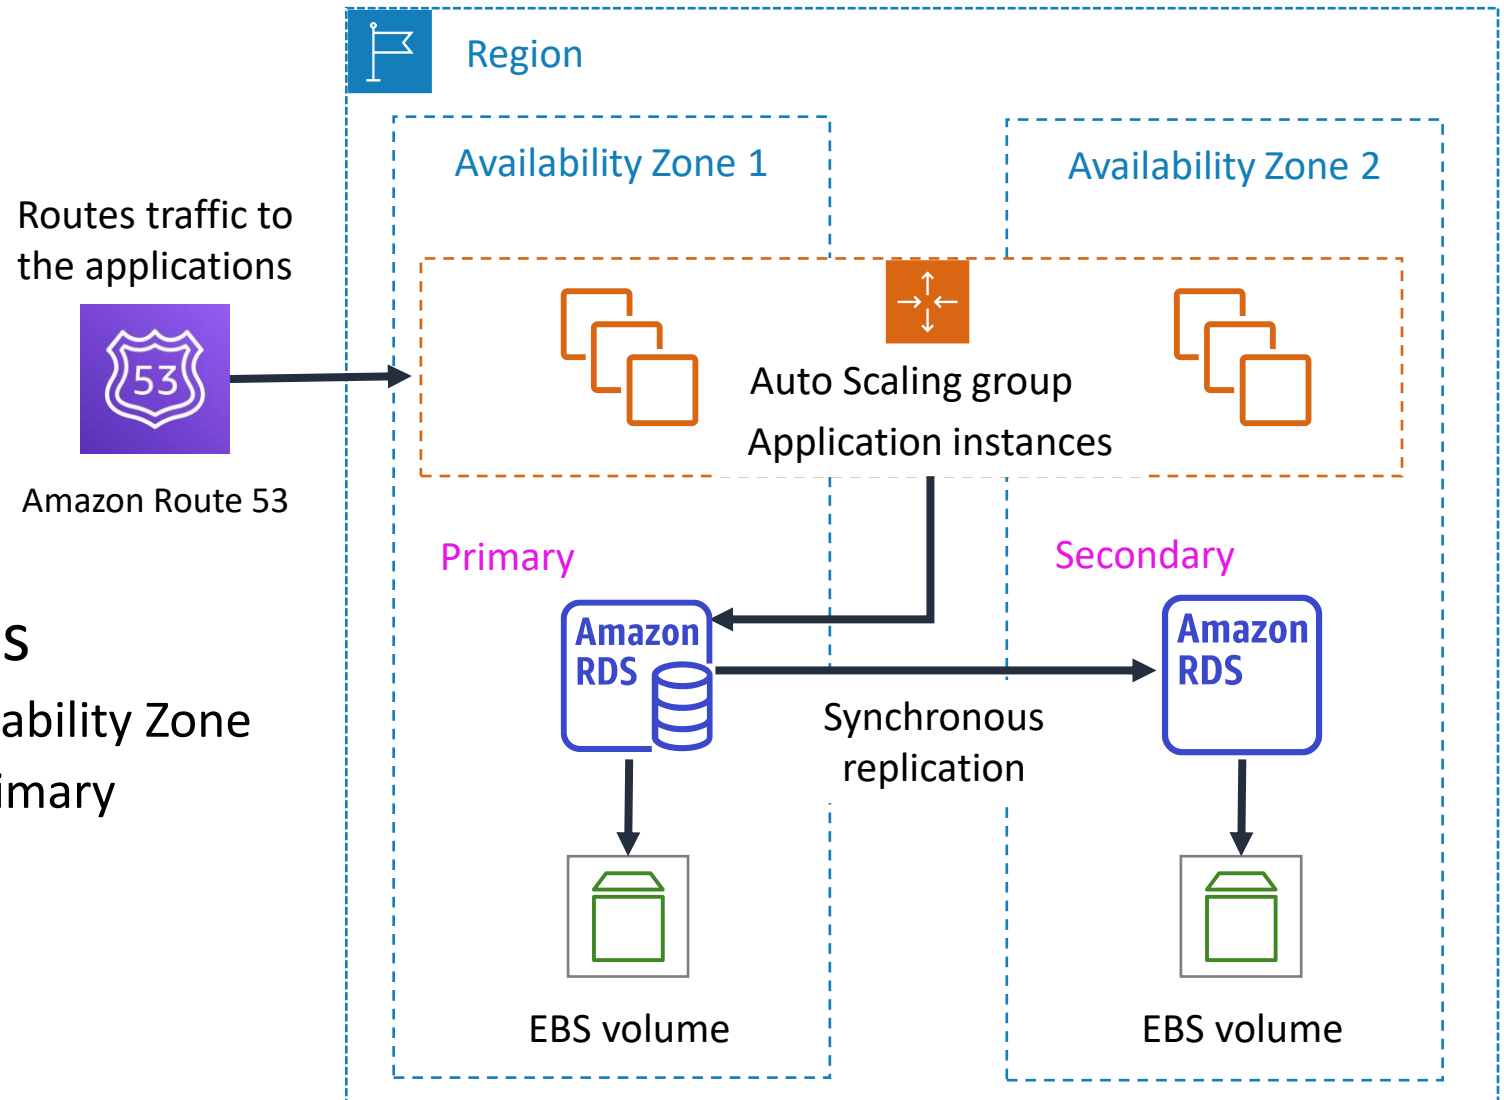
# Multi-AZ deployment for high availability

## Benefits

- Enhanced durability
- Increased availability
- Fail over to standby occurs automatically

## Automated failover conditions

- Loss of availability in primary Availability Zone
- Loss of network connectivity to primary
- Compute unit failure on primary
- Storage failure on primary



# Read replicas for performance

## Benefits

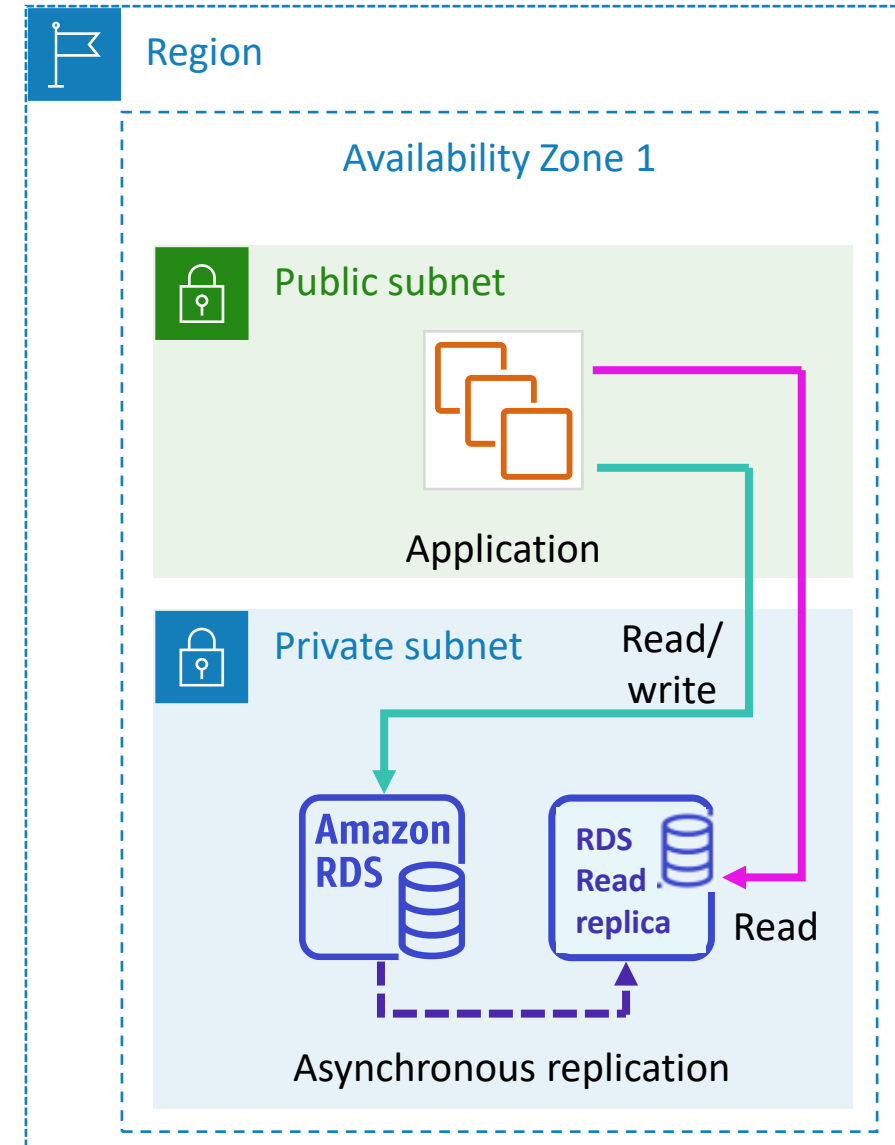
- Enhanced performance
- Increased availability
- Designed for security

## Supported by

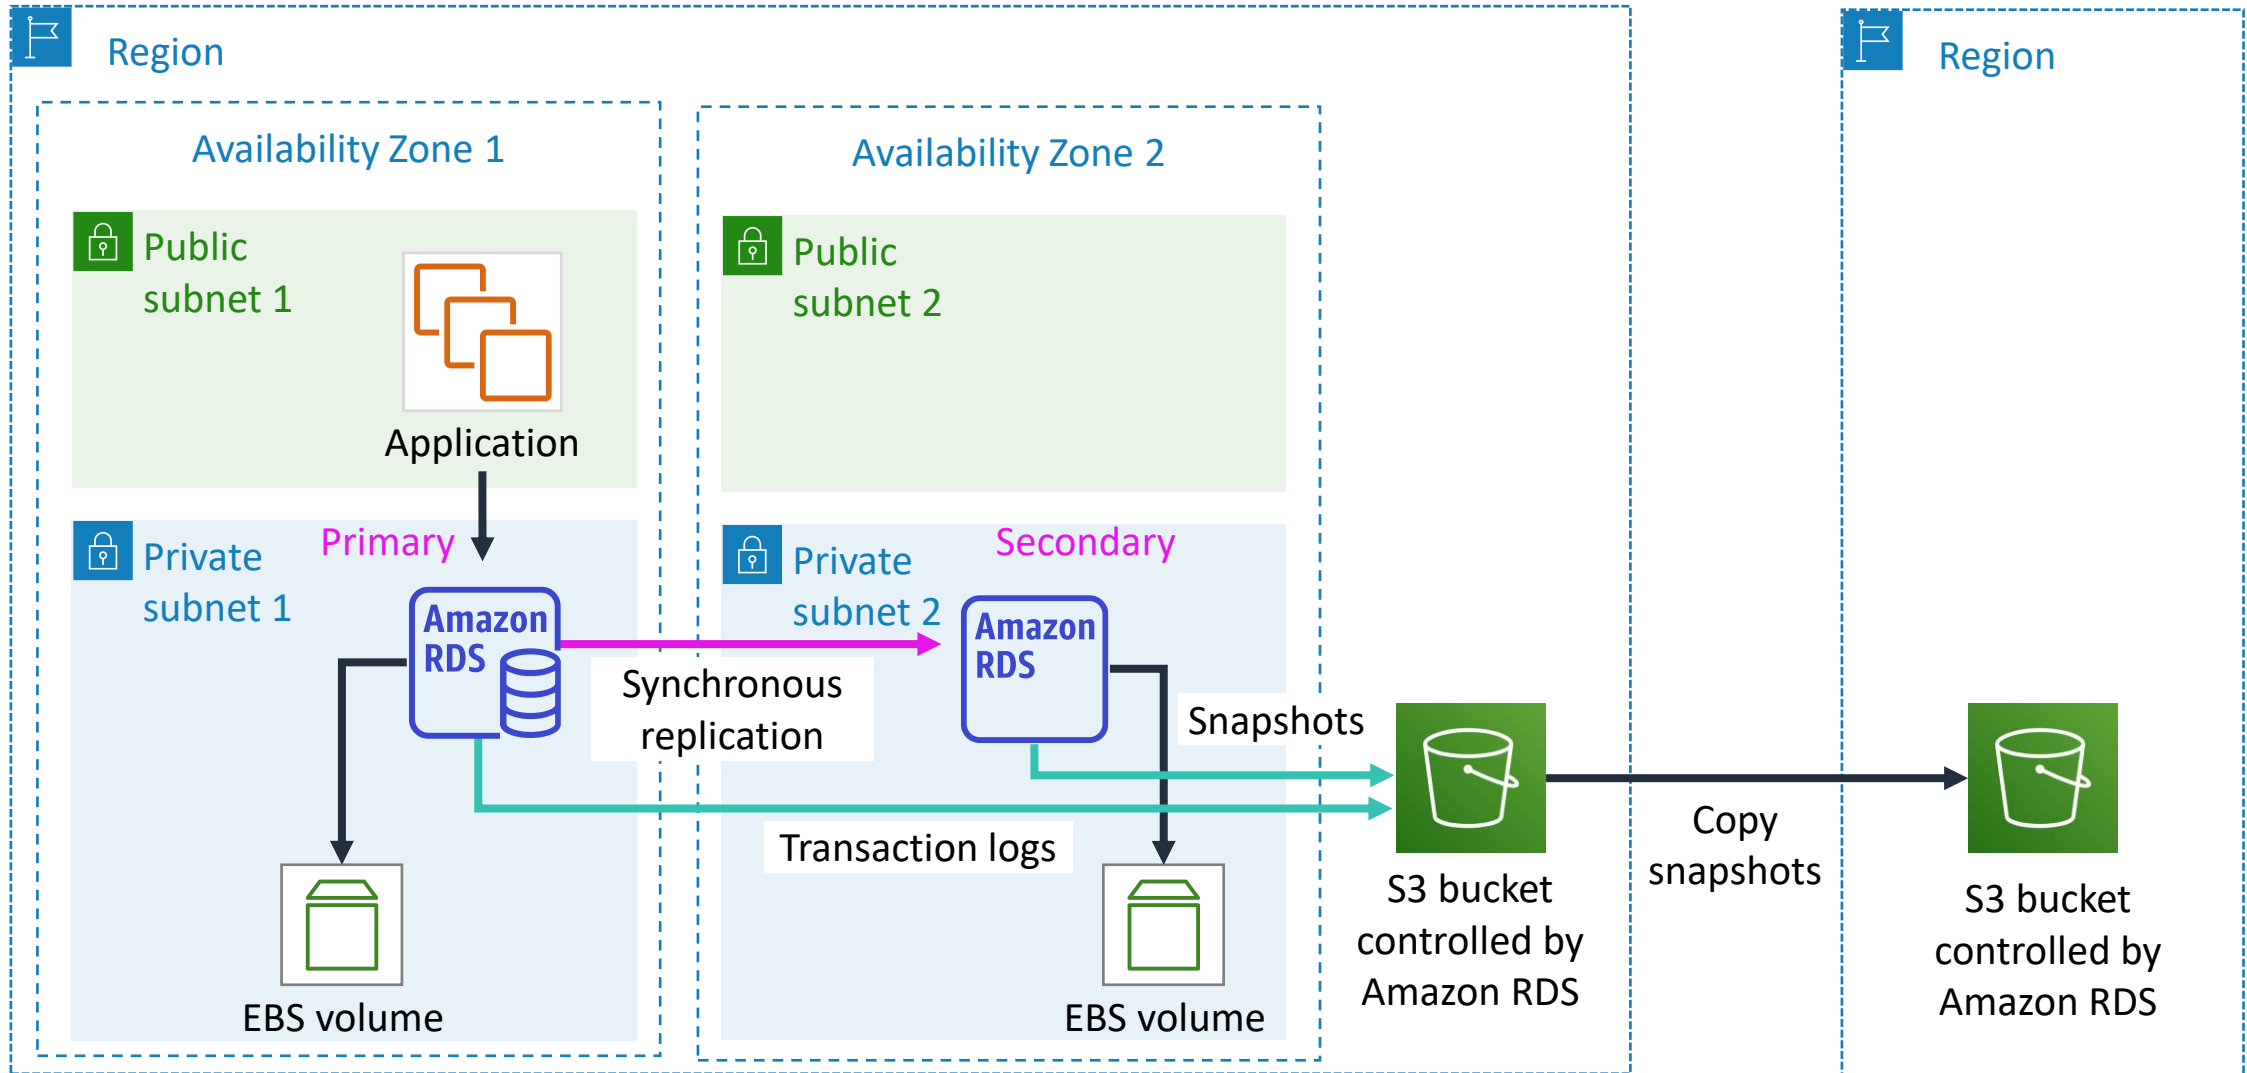
- MySQL
- MariaDB
- PostgreSQL
- Oracle

## Limits

- Five read replicas per primary
- For strict read-after-write consistency, read from the primary



# Amazon RDS backup solution





# Demonstration: Amazon RDS Automated Backup and Read Replicas





Amazon  
Aurora

Amazon Aurora is a fully managed, MySQL- and PostgreSQL-compatible, relational database engine.

- Used for online transactional processing (OLTP)
- Provides up to five times the throughput of MySQL\*
- Provides up to three times the throughput of PostgreSQL\*
- Replicates data six ways across three Availability Zones
- Requires little change to your existing application

\* Benchmarking details are available for [MySQL](#) and [PostgreSQL](#).





Amazon  
Redshift

Amazon Redshift is a **data warehousing** service.

- Is used for online analytics processing (OLAP)
- Stores very large datasets
  - Store highly structured, frequently accessed data in Amazon Redshift
  - Can also store exabytes of structured, semistructured, and unstructured data in Amazon S3

# Section 3 key takeaways



- **Managed AWS database services** handle administration tasks so you can focus on your applications
- Amazon RDS supports **Microsoft SQL Server, Oracle, MySQL, PostgreSQL, Aurora, and MariaDB**
- Amazon RDS **Multi-AZ deployments** provide **high availability** with automatic failover
- You can have up to five **read replicas** per primary database to improve Amazon RDS **performance**
- **Amazon Aurora** is a fully managed, MySQL- and PostgreSQL-compatible, relational database engine
- Amazon Redshift is a relational database offering for **data warehousing**

## Module 5: Adding a Database Layer


# Section 4: Amazon DynamoDB

# Amazon DynamoDB




Amazon  
DynamoDB

A fully managed **non-relational**  
**key-value** and **document** database service.



Performance  
at any scale

Extreme horizontal  
scaling capability



Serverless

Event-driven  
programming  
(serverless computing)



Enterprise-ready

Encryption, access  
controls, backups




Amazon  
DynamoDB

## Works well for applications that:

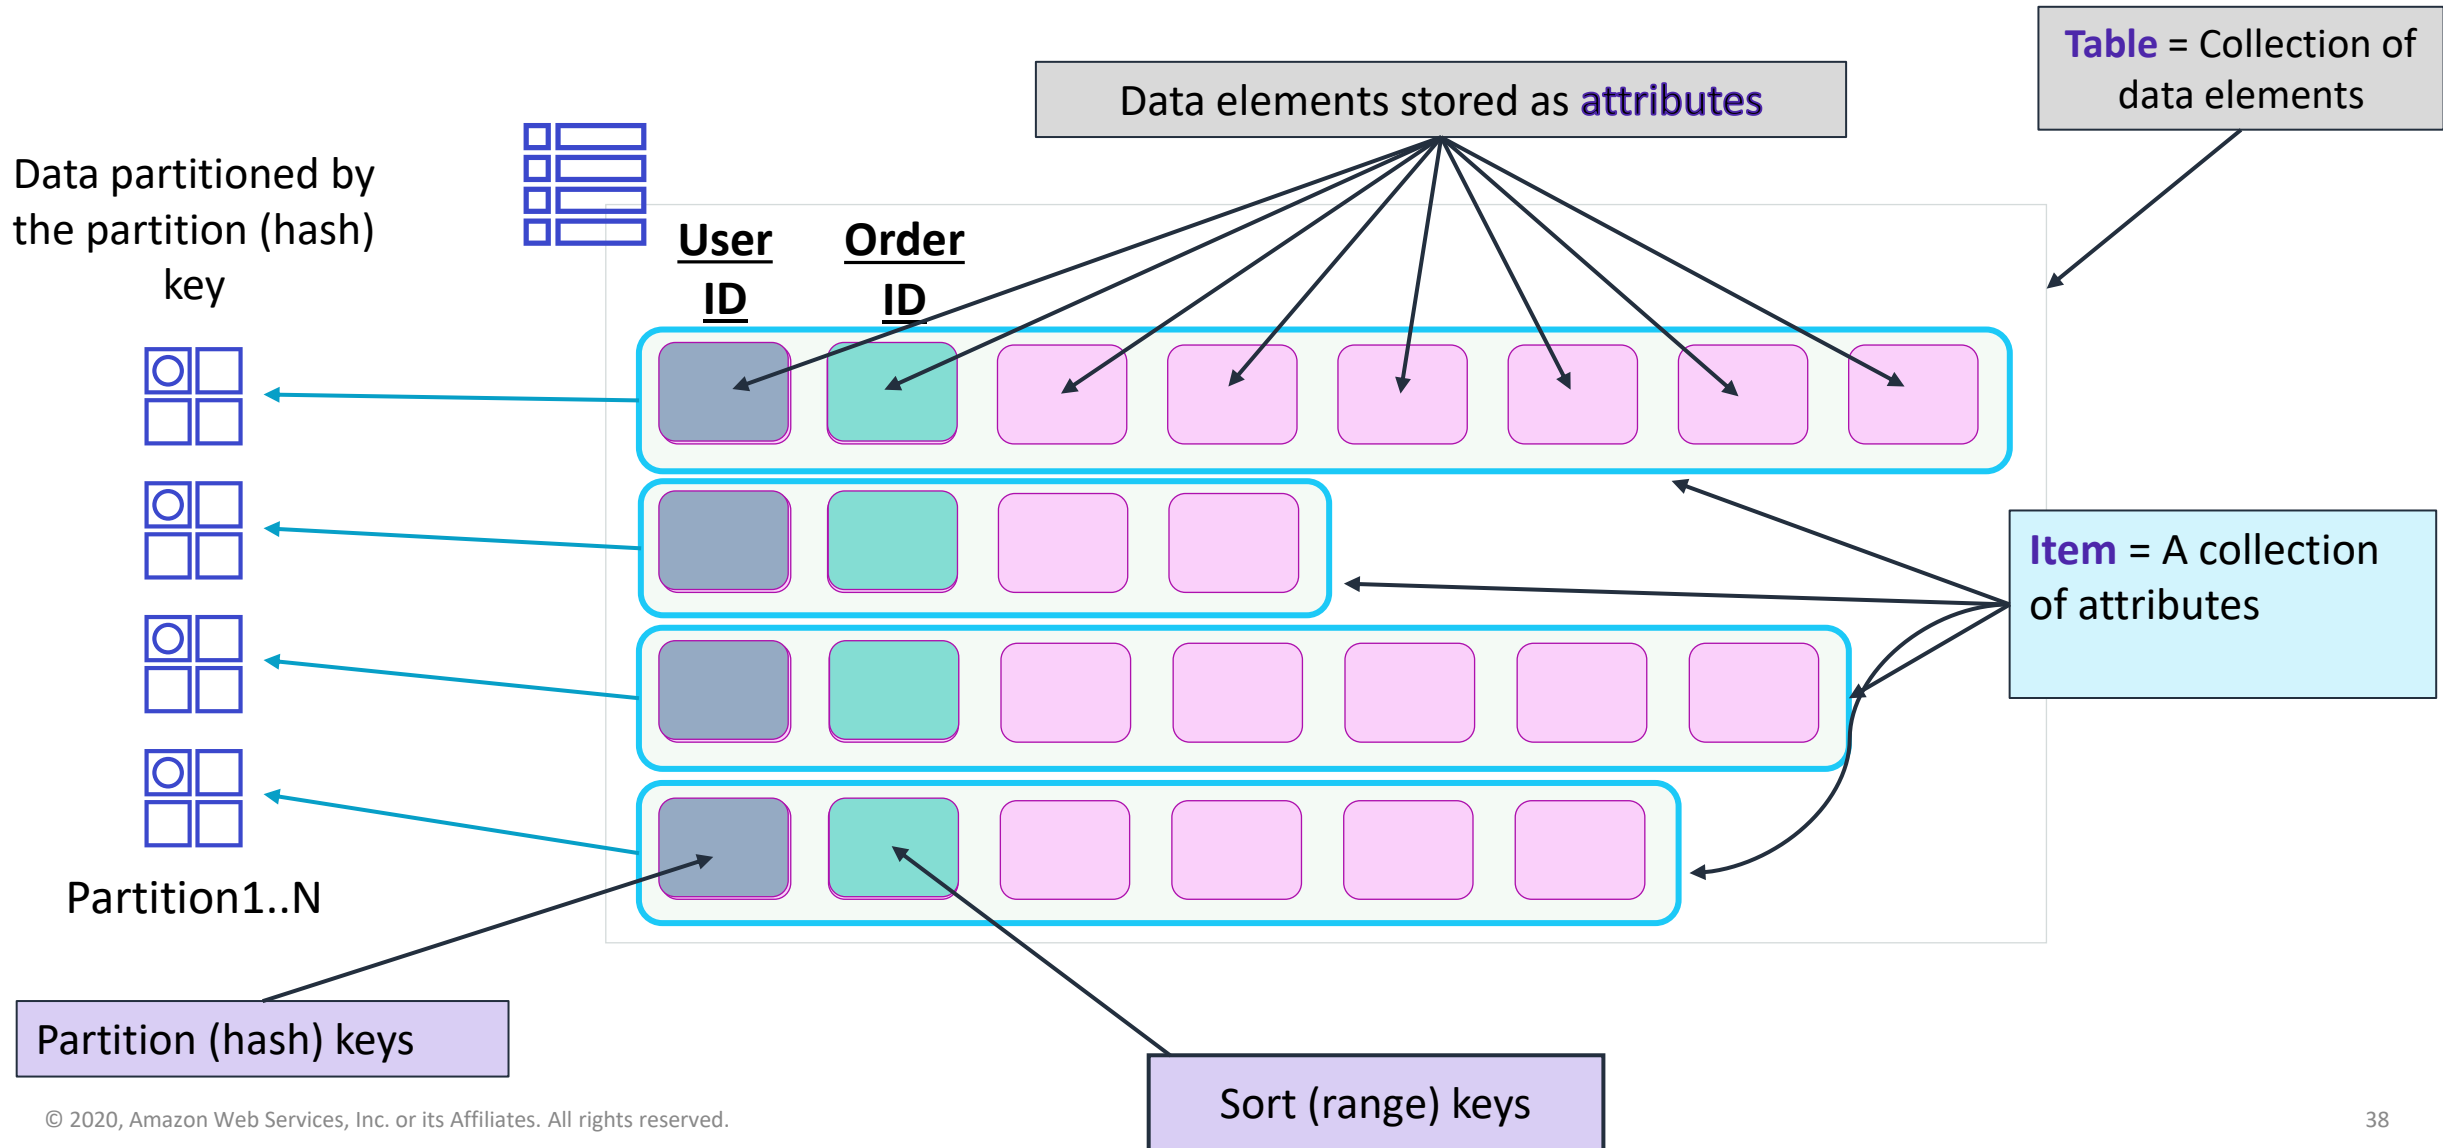
- Have simple high-volume data (high-TB range)
- Must scale quickly
- Don't need complex joins
- Require ultra-high throughput and low latency

## Key features

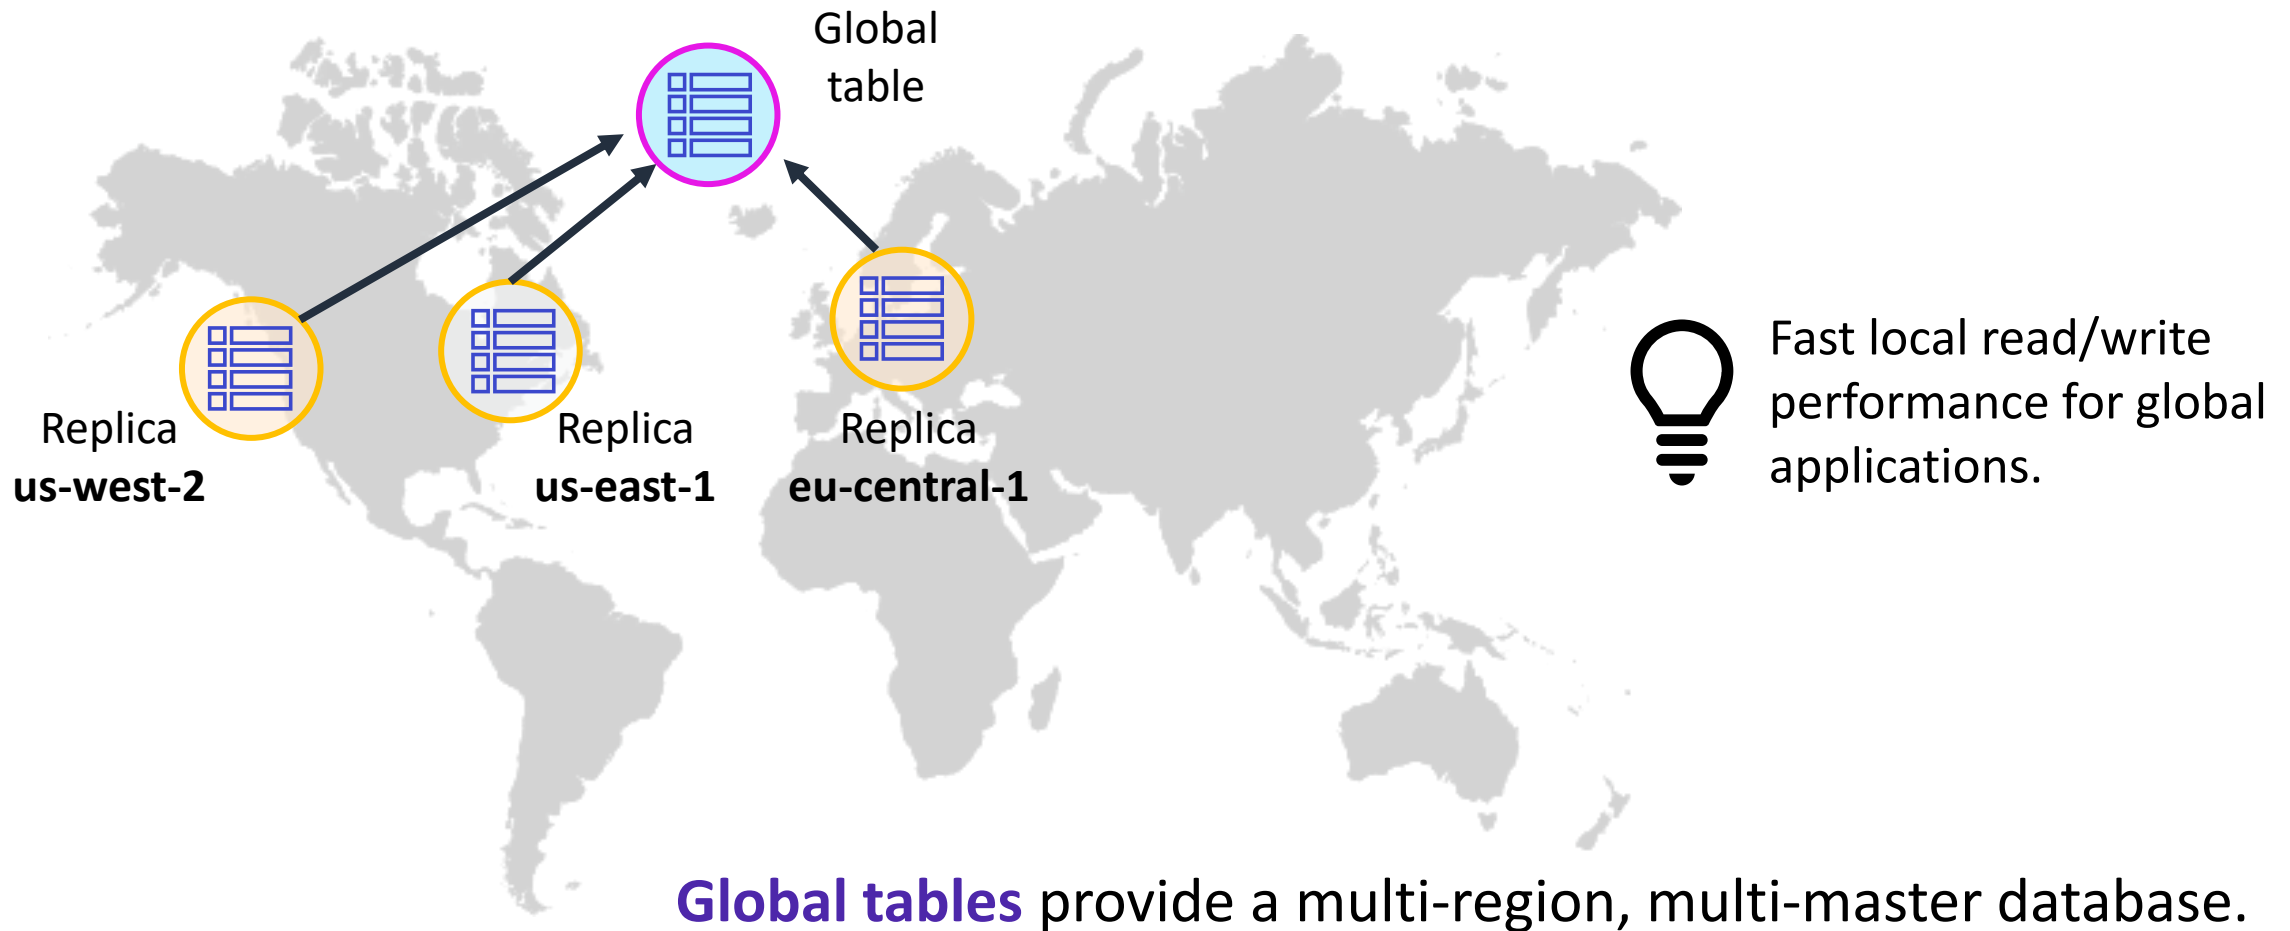
- NoSQL tables
- Items can have differing attributes
- In-memory caching
- Support for peaks of more than 20 million requests per second




# Amazon DynamoDB data model



# Amazon DynamoDB global tables



## Leaderboards and Scoring

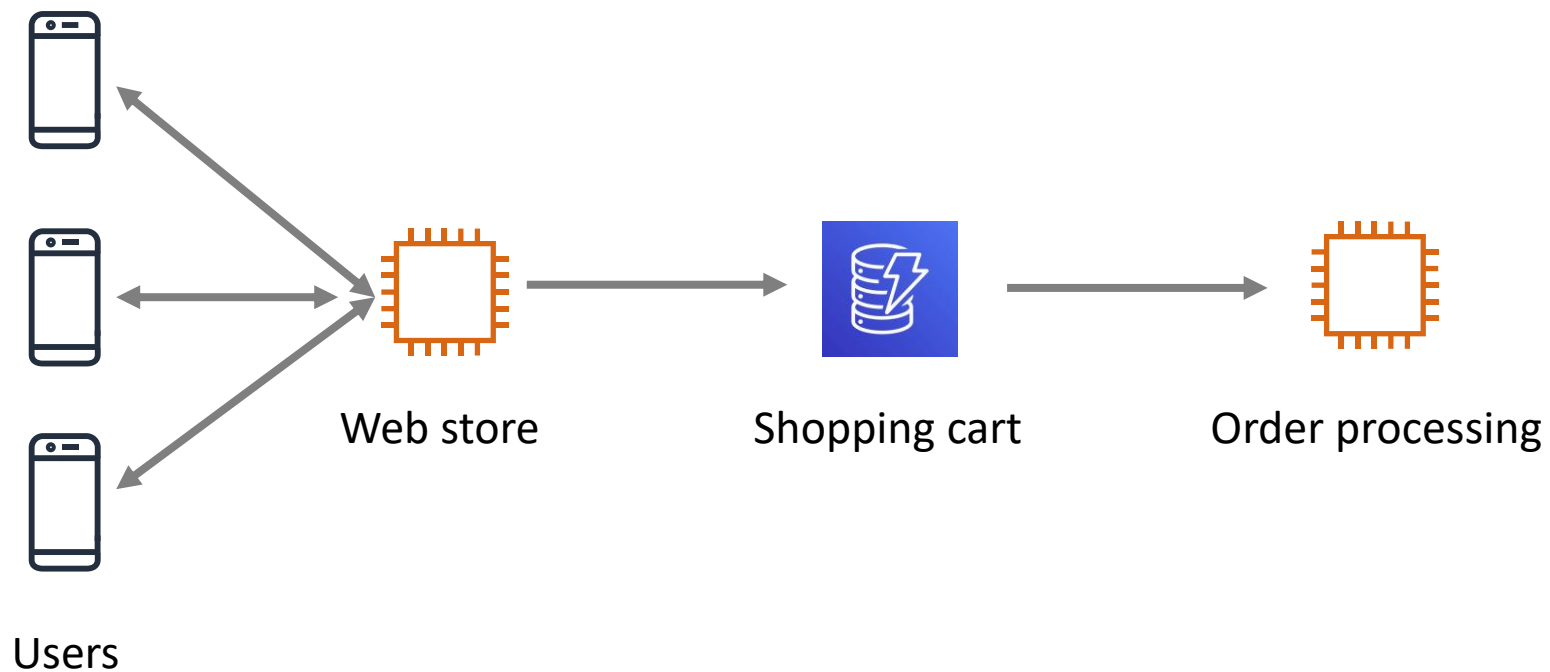


GameScores

UserId	GameTitle	TopScore	TopScoreDateTime	Wins	Losses	
"101"	"Galaxy Invaders"	5842	"2015-09-15:17:24:31"	21	72	...
"101"	"Meteor Blasters"	1000	"2015-10-22:23:18:01"	12	3	...
"101"	"Starship X"	24	"2015-08-31:13:14:21"	4	9	...
"102"	"Alien Adventure"	192	"2015-07-12:11:07:56"	32	192	...
"102"	"Galaxy Invaders"	0	"2015-09-18:07:33:42"	0	5	...
"103"	"Attack Ships"	3	"2015-10-19:01:13:24"	1	8	...
"103"	"Galaxy Invaders"	2317	"2015-09-11:06:53:00"	40	3	...
"103"	"Meteor Blasters"	723	"2015-10-19:01:13:24"	22	12	...
"103"	"Starship X"	42	"2015-07-11:06:53:00"	4	19	...
...	...	...	...	...	...	...




## Temporary Data (Online Cart)




# Amazon DynamoDB consistency options

## Eventually consistent




The **default** setting. All copies of data usually reach consistency within **1 second**.

## Strongly consistent




This feature is optional. Use for applications that require all reads to return a result that reflects all writes before the read.

# Class discussion: Which database should the café use?



Amazon  
RDS



Amazon  
DynamoDB



## Section 4 key takeaways



- Amazon DynamoDB is a fully managed non-relational **key-value** and document **NoSQL** database service.
- DynamoDB is **serverless**, provides extreme **horizontal scaling** and **low latency**.
- DynamoDB **global tables** ensure that data is replicated to multiple Regions.
- DynamoDB provides **eventual consistency** by default (in general, it is fully consistent for reads 1 second after the write). **Strong consistency** is also an option.

## Module 5: Adding a Database Layer

# Section 5: Database security controls

# Securing Amazon RDS databases

## Recommendations

- Run the RDS instance in a [virtual private cloud \(VPC\)](#)
  - Provides service isolation and IP firewall protection
- Use [AWS Identity and Access Management \(IAM\) policies](#) for authentication and access
  - Permissions determine who is allowed to manage Amazon RDS resources
- Use [security groups](#) to control what IP addresses or EC2 instances can connect to your databases
  - By default, network access is disabled
- Use [Secure Sockets Layer \(SSL\)](#) for encryption in transit
- Use Amazon RDS [encryption](#) on DB instances and snapshots to secure data at rest
- Use the [security features of your DB engine](#) to control who can log in to the databases on a DB instance
- Configure event notifications to alert you when important Amazon RDS events occur




Amazon RDS



## Recommendations

- Use **IAM roles** to authenticate access
- Use **IAM policies**
  - To define fine-grain access permissions to use DynamoDB APIs
  - Define access at the table, item, or attribute level
  - Follow the **principle of granting least privilege**
- Configure **VPC endpoints**
  - Prevents connection traffic from traversing the open internet
  - VPC endpoint policies allow you to control and limit API access to a DynamoDB table
- Consider **client-side encryption**
  - Encrypt data as close as possible to its origin



Amazon  
DynamoDB

## Security provided by default

- **Encryption at rest** of all user data stored in tables, indexes, streams, and backups
- **Encryption in transit** – All communications to and from DynamoDB and other AWS resources use HTTPS

## Module 5: Adding a Database Layer

# Section 6: Migrating data into AWS databases





# AWS Database Migration Service



AWS Database Migration Service (AWS DMS)

- Use to migrate to and from most commercial and open source databases
- Migrate between databases on Amazon EC2, Amazon RDS, Amazon S3, and on-premises

Example  
homogenous  
migration



Example  
heterogeneous  
migration

# AWS DMS key features

- Perform one-time migrations
- Or, accomplish continuous data replication
  - Example: Configure continuous data replication of an on-premises database to an RDS instance
- **AWS Schema Conversion Tool (AWS SCT)** supports changing the database engine between source and target
- Typical migration major steps:
  1. Create a target database
  2. Migrate the database schema
  3. Set up the data replication process
  4. Initiate the data transfer, and confirm completion
  5. Switch production to the new database (for one-time migrations)



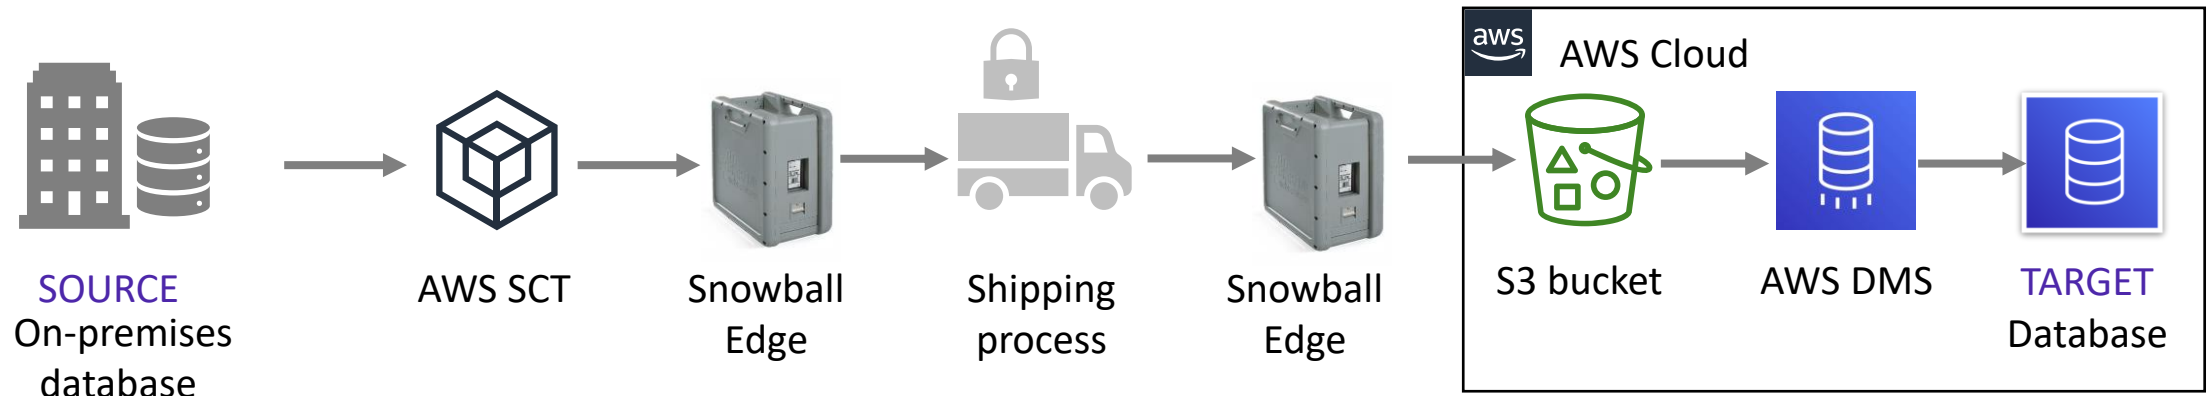
# Using AWS Snowball Edge with AWS DMS

When migrating data is not practical:


- Database is too large
- Connection is too slow
- You have privacy and security concerns

## Use **AWS Snowball Edge**

- Multi-terabyte transfer without using the internet




## Module 5 - Challenge Lab: Migrating a Database to Amazon RDS





# The business need: A managed database

The database that runs on the EC2 instance is becoming difficult for Sofía and Nikhil to maintain.




When Olivia visited the café recently, she told them about the features of Amazon RDS.

Sofía and Nikhil decided to migrate the café's database to Amazon RDS.

# Challenge lab: Customer requirements

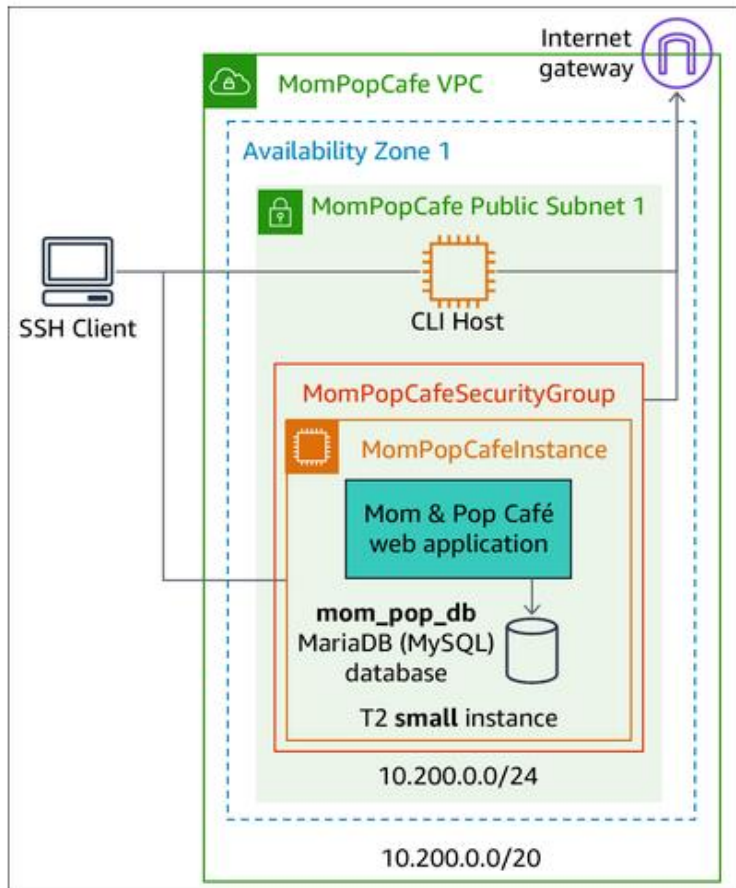
1. Create an Amazon RDS instance
2. Migrate application data to the Amazon RDS instance
3. Configure the website to use the Amazon RDS instance
4. Monitor the Amazon RDS database



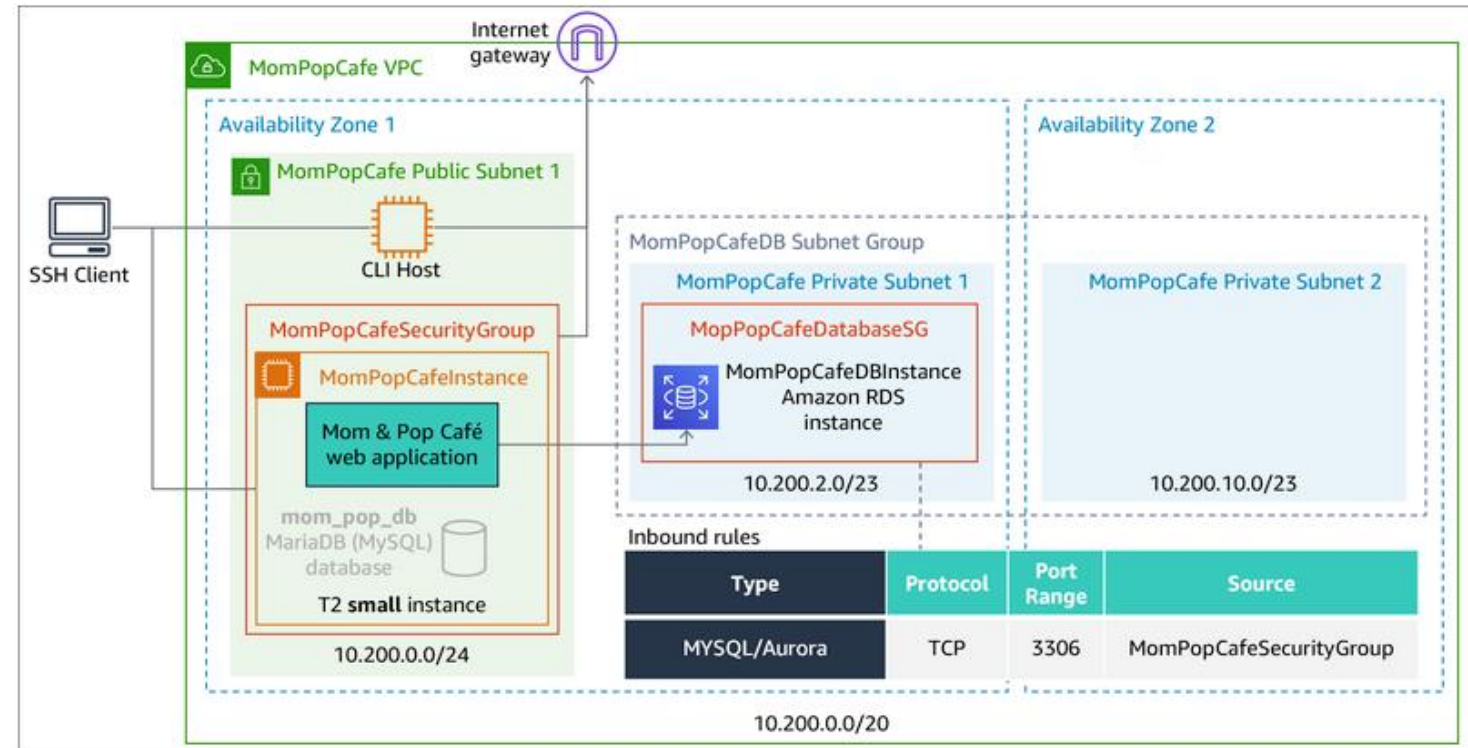
# Challenge lab: Final product

## PLACEHOLDER


### Before migration



### After migration





A top-down photograph of a teal ceramic coffee cup filled with dark coffee, topped with a layer of white foam. The cup sits on a matching teal saucer. To the left of the cup is a wooden scoop filled with dark brown coffee beans, with several beans scattered on the dark, textured surface around it. A light-colored, textured cloth is partially visible on the left side of the frame.

## Begin Module 5 – Challenge Lab: Migrating a Database to Amazon RDS



# Challenge lab debrief: Key takeaways



## Module 5: Adding a Database Layer

# Module wrap-up

In summary, in this module, you learned how to:

- Compare database types
- Differentiate between managed versus unmanaged services
- Explain when to use Amazon Relational Database Service (Amazon RDS)
- Explain when to use Amazon DynamoDB
- Describe available database security controls
- Describe how to migrate data into Amazon Web Services (AWS) databases
- Deploy a database server

# Complete the knowledge check





# Sample exam question

An application requires a highly available relational database with an initial storage capacity of 8 TB. The database will grow by 8 GB every day. To support expected traffic, at least eight read replicas will be required to handle database reads.

Which option will meet these requirements?

- A. DynamoDB
- B. Amazon S3
- C. Amazon Aurora
- D. Amazon Redshift

# Additional resources

- [AWS Databases – Resource page](#)
- [Amazon RDS Getting Started Guide](#)
- [Best Practices for Amazon RDS](#)
- [Amazon RDS FAQs](#)
- [Amazon DynamoDB Developer Guide](#)
- [Amazon DynamoDB FAQs](#)



# Thank you

© 2020 Amazon Web Services, Inc. or its affiliates. All rights reserved. This work may not be reproduced or redistributed, in whole or in part, without prior written permission from Amazon Web Services, Inc. Commercial copying, lending, or selling is prohibited. Corrections or feedback on the course, please email us at: [aws-course-feedback@amazon.com](mailto:aws-course-feedback@amazon.com). For all other questions, contact us at: <https://aws.amazon.com/contact-us/aws-training/>. All trademarks are the property of their owners.

